

In [1]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files
under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved
as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of
the current session

/kaggle/input/30-days-of-ml/sample_submission.csv
/kaggle/input/30-days-of-ml/train.csv
/kaggle/input/30-days-of-ml/test.csv
```

Loading All Datasets

In [2]:

```
# Use the column 'id' as the index for train and test datasets
train = pd.read_csv('/kaggle/input/30-days-of-ml/train.csv', index_col = 'id')
test = pd.read_csv('/kaggle/input/30-days-of-ml/test.csv', index_col = 'id')
sample = pd.read_csv('/kaggle/input/30-days-of-ml/sample_submission.csv')
```

In [3]:

```
# View train df
print(train.head())
```

	cat0	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	...	cont5	cont6	\
id											...			
1	B	B	B	C	B	B	A	E	C	N	...	0.400361	0.160266	
2	B	B	A	A	B	D	A	F	A	O	...	0.533087	0.558922	
3	A	A	A	C	B	D	A	D	A	F	...	0.650609	0.375348	
4	B	B	A	C	B	D	A	E	C	K	...	0.668980	0.239061	
6	A	A	A	C	B	D	A	E	A	N	...	0.686964	0.420667	

	cont7	cont8	cont9	cont10	cont11	cont12	cont13	\
id								
1	0.310921	0.389470	0.267559	0.237281	0.377873	0.322401	0.869850	
2	0.516294	0.594928	0.341439	0.906013	0.921701	0.261975	0.465083	
3	0.902567	0.555205	0.843531	0.748809	0.620126	0.541474	0.763846	
4	0.732948	0.679618	0.574844	0.346010	0.714610	0.540150	0.280682	
6	0.648182	0.684501	0.956692	1.000773	0.776742	0.625849	0.250823	

	target
id	
1	8.113634
2	8.481233
3	8.364351
4	8.049253
6	7.972260

[5 rows x 25columns]

In [4]:

```
# View test df
print(test.head())
```

	cat0	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	...	cont4	cont5	\
id											...			
0	B	B	B	C	B	B	A	E	E	I	...	0.476739	0.376350	
5	A	B	A	C	B	C	A	E	C	H	...	0.285509	0.860046	
15	B	A	A	A	B	B	A	E	D	K	...	0.697272	0.683600	
16	B	B	A	C	B	D	A	E	A	N	...	0.719306	0.777890	
17	B	B	A	C	B	C	A	E	C	F	...	0.313032	0.431007	

	cont6	cont7	cont8	cont9	cont10	cont11	cont12	\
id								
0	0.337884	0.321832	0.445212	0.290258	0.244476	0.087914	0.301831	
5	0.798712	0.835961	0.391657	0.288276	0.549568	0.905097	0.850684	
15	0.404089	0.879379	0.275549	0.427871	0.491667	0.384315	0.376689	
16	0.730954	0.644315	1.024017	0.391090	0.988340	0.411828	0.393585	
17	0.390992	0.408874	0.447887	0.390253	0.648932	0.385935	0.370401	

	cont13
id	
0	0.845702
5	0.693940
15	0.508099
16	0.461372
17	0.900412

[5 rows x 24 columns]

In [5]:

```
# View sample submission
print(sample.head())
```

	id	target
0	0	0.5
1	5	0.5
2	15	0.5
3	16	0.5
4	17	0.5

In [6]:

```
# Separate target and features
y = train['target']
features = train.drop(['target'], axis = 1)
```

Ordinal Encode Categorical Columns In Given Features

In [7]:

```
from sklearn.preprocessing import OrdinalEncoder

object_cols = [col for col in features.columns if 'cat' in col]

X = features.copy()
X_test = test.copy()

# ordinal-encode the categorical columns in X and X_test datasets.
ordinal_encoder = OrdinalEncoder()
X[object_cols] = ordinal_encoder.fit_transform(features[object_cols])
X_test[object_cols] = ordinal_encoder.transform(test[object_cols])
```

In [8]:

```
# Preview
print(X.head())
```

	cat0	cat1	cat2	cat3	cat4	cat5	cat6	cat7	cat8	cat9	...	cont4	\
id											...		
1	1.0	1.0	1.0	2.0	1.0	1.0	0.0	4.0	2.0	13.0	...	0.610706	
2	1.0	1.0	0.0	0.0	1.0	3.0	0.0	5.0	0.0	14.0	...	0.276853	
3	0.0	0.0	0.0	2.0	1.0	3.0	0.0	3.0	0.0	5.0	...	0.285074	
4	1.0	1.0	0.0	2.0	1.0	3.0	0.0	4.0	2.0	10.0	...	0.284667	
6	0.0	0.0	0.0	2.0	1.0	3.0	0.0	4.0	0.0	13.0	...	0.287595	

	cont5	cont6	cont7	cont8	cont9	cont10	cont11	\
id								
1	0.400361	0.160266	0.310921	0.389470	0.267559	0.237281	0.377873	
2	0.533087	0.558922	0.516294	0.594928	0.341439	0.906013	0.921701	
3	0.650609	0.375348	0.902567	0.555205	0.843531	0.748809	0.620126	
4	0.668980	0.239061	0.732948	0.679618	0.574844	0.346010	0.714610	
6	0.686964	0.420667	0.648182	0.684501	0.956692	1.000773	0.776742	

	cont12	cont13
id		
1	0.322401	0.869850
2	0.261975	0.465083
3	0.541474	0.763846
4	0.540150	0.280682
6	0.625849	0.250823

[5 rows x 24 columns]

In [9]:

```
# Split the encoded features df into training and validation datasets
from sklearn.model_selection import train_test_split

X_train, X_valid, y_train, y_valid = train_test_split(X, y, random_state = 0)
```

Use A XGBRegressor Model And Train It To Provide Better Performance

In [10]:

```
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error
```

In [11]:

```
# Define the model
model = XGBRegressor(learning_rate = 0.02, random_state = 42, tree_method = 'gpu_hist')

# Train the model(will take some time to execute)
model.fit(X_train, y_train, early_stopping_rounds = 5, eval_set = [(X_valid, y_valid)],
verbose = False)
preds_valid = model.predict(X_valid)
print(mean_squared_error(y_valid, preds_valid, squared = False))

1.262424849715232
```

In [12]:

```
# Get the first 5 prediction results
print(list(preds_valid)[0:6:1])

[7.189719, 7.0716166, 7.257436, 7.1754956, 7.313472, 7.2468147]
```

Submission

In [13]:

```
# Use the model to generate predictions on the given testing dataset
test_predictions = model.predict(X_test)
```

```
# Save the predictions to a CSV file
output = pd.DataFrame({'Id': X_test.index, 'target': test_predictions})
output.to_csv('submission.csv', index = False)
```

In [14]:

```
# Preview predicted output
print(output.head(10))
```

	Id	target
0	0	7.033919
1	5	7.255158
2	15	7.293473
3	16	7.254006
4	17	7.162273
5	19	7.237772
6	20	7.246070
7	21	7.082504
8	23	7.296534
9	29	7.266431